

# CMPUT-379 Lab 1

---

Akemi Izuko, Armaan Katyal, Ellis McDougald, Han Yang, Patrick Zijlstra, Shasta Johnsen-Sollos, Steven Oufan Hai, Zhaoyu Li

# Today's lab

- Get familiar with the programming environment
  - SSH and Git
- Basic steps for Unix system programming
  - Makefiles
  - GDB and Valgrind
  - Man pages
  - Utilities and commands
    - `objdump`, `nm`, `strace`, `ps`, `pstree`, `top`, `watch`, `kill`
  - Debugging:
    - `errno` & `perror`
- System calls for process management
  - `fork()`, `execve()`, `getpid()`, `wait()`, `waitpid()`, `exit()`

# Getting started with the programming environment

# CS Linux Machines

- Make sure your programs can compile and run on these machines
- You must connect to the ualberta vpn to access these, or proxyjump through [login.cs.ualberta.ca](https://login.cs.ualberta.ca) (advanced)
- **Assignments will be graded on these machines**

## Lab machines:

[ucomm-2030-wXX.cs.ualberta.ca](https://ucomm-2030-wXX.cs.ualberta.ca)  
(XX must be between 01 and 04 inclusive)  
[ucomm-2070-wXX.cs.ualberta.ca](https://ucomm-2070-wXX.cs.ualberta.ca)  
(XX must be between 00 and 24 inclusive)  
[ucomm-2086-wXX.cs.ualberta.ca](https://ucomm-2086-wXX.cs.ualberta.ca)  
(XX must be between 00 and 33 inclusive)  
[ucomm-2130-wXX.cs.ualberta.ca](https://ucomm-2130-wXX.cs.ualberta.ca)  
(XX must be between 00 and 25 inclusive)  
[ucomm-2140-wXX.cs.ualberta.ca](https://ucomm-2140-wXX.cs.ualberta.ca)  
(XX must be between 00 and 25 inclusive)  
[ucomm-3130-wXX.cs.ualberta.ca](https://ucomm-3130-wXX.cs.ualberta.ca)  
(XX must be between 00 and 23 inclusive)  
[ucomm-3140-wXX.cs.ualberta.ca](https://ucomm-3140-wXX.cs.ualberta.ca)  
(XX must be between 00 and 21 inclusive)

# Connect to the CS Linux machines with **ssh**

- Open your terminal app
- Run **ssh** CCID@ucomm-2070-wXX.cs.ualberta.ca
- Enter your password for you CCID
- **exit** or Ctrl-D to exit
- For Windows you can use PuTTY/MobaXTerm/WSL terminal

```
jihoon@Phoenix ➤ ssh og@ug00.cs.ualberta.ca
og@ug00.cs.ualberta.ca's password:
Welcome to Ubuntu 20.04.6 LTS (GNU/Linux 5.15.0-119-generic x86_64)

Department of Computing Science
University of Alberta


Unauthorized use is prohibited.

Problem reports can be made using mail to ist@ualberta.ca
or https://www.ualberta.ca/computing-science/links-and-resources/technical-support

Last login: Sat Sep 14 19:02:18 2024 from 162.157.230.82
og@ug00:~>
```

# Editing files on remote machines

- Terminal text editors like **vim**, **emacs**, **nano**
- VSCode Remote Development
  - Modern GUI application with commonly used keybinds
  - Runs on your own machine
  - Manages and edit your remote files through SSH



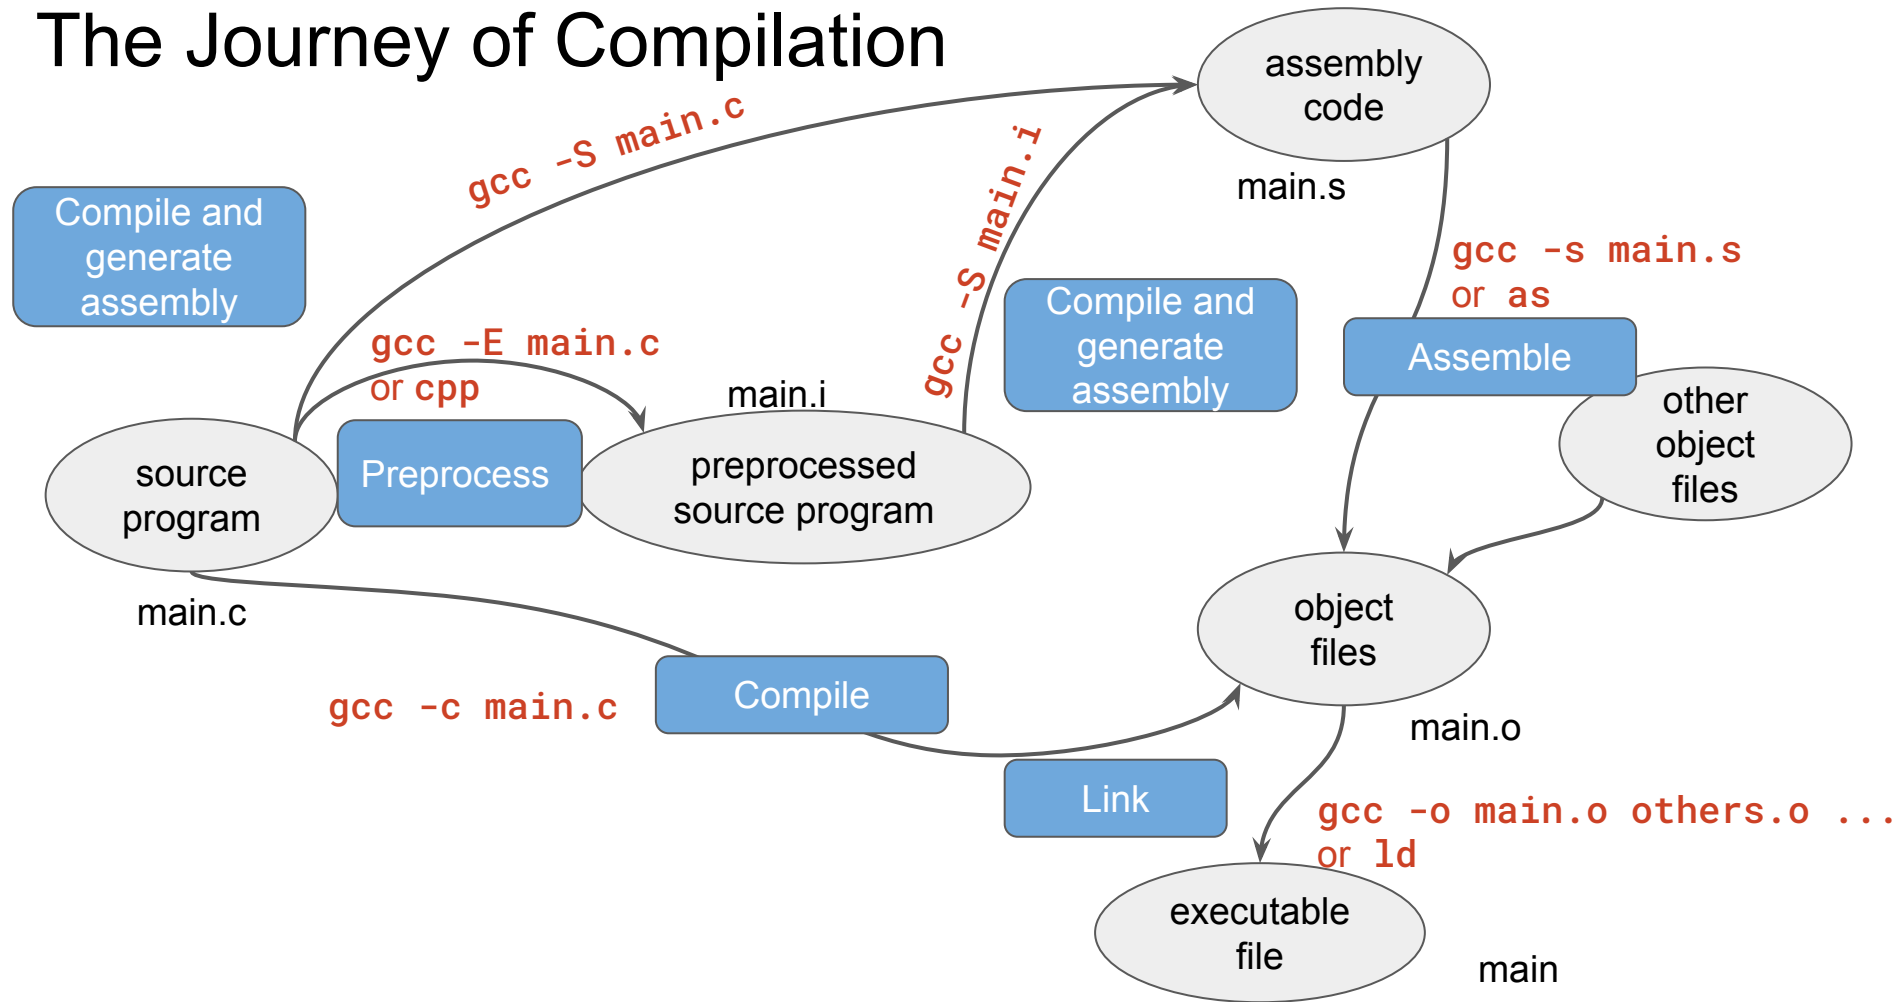
# Quick git tutorial

- Cloning a repository onto your local machine
  - `git clone <url>`
- Adding files to the staging area
  - `git add <filename>`
- Check status of files in the local repository
  - `git status`
- Commit files in your staging area to your local git repository
  - `git commit -m <short message>`
- Pushing/Pulling commits to and from GitHub
  - `git pull`
  - `git push`


# UNIX Programming 101



# The Journey of Compilation



# The Simplified Journey of Compilation



# Compilation options

- Use command line options to control the behaviour of **gcc**
- **-o <output filename>** output file name (create executable if **-c** is not specified)
- **-c [output filename]** create an object file (if no output filename then it will use the source code filename appended with **.o**)
- **-g** keep debugging information
- **-Wall** adds most warnings
  - Your assignment **must** compile cleanly (as in no warnings) for full marks

# Compilation - compile and link

- Some files do not need to be recompiled
- Object files can be reused/shared
- Use **make** to help you automate this process - covered next

# Make - Introduction

- **Make**: a tool to automate compilation
- **REQUIRED FOR ASSIGNMENTS**
- When properly setup it should only recompile outdated files
- You will need a **Makefile** in your project folder

# make - Makefile basics

- Hello world in Makefile

```
say_hello:  
    echo "Hello World"
```

- Run it in shell

```
$ make  
echo "Hello World"  
Hello World
```

# make - Makefile basics

- Syntax to define rules

```
target: prerequisites  
<TAB> recipe
```

- Run it in shell

```
$ make <target>
```

# make - Makefile basics

- When we run `make <target>` in the shell, `make` will
  - Check the dependencies of the target
    - If any of the dependencies have been modified since the last time the `<target>` was generated it will
      - Run the recipe line by line
- You will need at least 3 targets in your Makefile for your assignment
  - compile
  - link
  - clean



# make - Makefile basics

- Dependencies can be another rule's target!
- Putting them all together

```
code_piece1.o: code_piece1.c
    gcc -c code_piece1.c -o code_piece1.o

code_piece2.o: code_piece2.c
    gcc -c code_piece2.c -o code_piece2.o

awesome_app: code_piece1.o code_piece2.o
    gcc -o awesome_app code_piece1.o code_piece2.o

clean:
    rm *.o awesome_app
```

# make - Variables

```
CC      = gcc
CFLAGS  = -Wall
OBJECTS = code_piece1.o code_piece2.o
```

Define and assign variables

```
code_piece1.o: code_piece1.c
    $(CC) $(CFLAGS) -c code_piece1.c -o code_piece1.o

code_piece2.o: code_piece2.c
    $(CC) $(CFLAGS) -c code_piece2.c -o code_piece2.o

awesome_app: $(OBJECTS)
    $(CC) -o awesome_app $(OBJECTS)
```

Use  
variables

```
gcc -Wall -c code_piece1.c -o code_piece1.o
```

Equivalent to

# GDB

- A tool to debug your program
- Use it to find errors that's hard to address
- Need to add the flag **-g** when you compile your program

```
$(GG) $(CFLAGS) prog.c -o output -g
```

- Then, use gdb by:

```
gdb ./output
```

- Running **list** in gdb will display the code
- **break** <line#> will add a breakpoint at the specified line number
- **run** will execute the program from the start to finish or until the first breakpoint
- Important commands
  - **run** / **continue** / **next** / **step** / **until** / **print** / **call** / **quit** / **break** + **line #** / etc...

# Valgrind

- A tool used to check memory leaks within your program
- Like gdb requires the `-g` flag during compilation to get the line number(s) of where the problem originates

```
valgrind --leak-check=yes ./your_prog arg1 arg2 ...
```

## Valgrind example - leaky\_prog.c

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    int *p = malloc(64);
}
```

# Valgrind example output

```
==2747425==      in use at exit: 64 bytes in 1 blocks
==2747425==    total heap usage: 2 allocs, 1 frees, 1,088 bytes allocated
==2747425==
==2747425== 64 bytes in 1 blocks are definitely lost in loss record 1 of 1
==2747425==    at 0x483B7F3: malloc (in /usr/lib/x86_64-linux-gnu/valgrind/vgpreload_memcheck-amd64-linux.so)
==2747425==    by 0x109185: main (leaky_prog.c:5)
==2747425==
==2747425== LEAK SUMMARY:
==2747425==    definitely lost: 64 bytes in 1 blocks
==2747425==    indirectly lost: 0 bytes in 0 blocks
==2747425==    possibly lost: 0 bytes in 0 blocks
==2747425==    still reachable: 0 bytes in 0 blocks
==2747425==    suppressed: 0 bytes in 0 blocks
==2747425==
==2747425== Use --track-origins=yes to see where uninitialised values come from
==2747425== For lists of detected and suppressed errors, rerun with: -s
```

# Man page

- AKA manual page are a set of software documents for user commands, syscalls, and libraries for Unix and Unix-like OSes
- Usage: `man [section] name`
- Sections
  - 1 - user commands (e.g., `man man` - “What is man and how to use it”)
  - 2 - system calls (e.g., `man 2 fork` - “How to create a child process”)
  - 3 - C standard library (e.g., `man 3 printf` - “How do I print in C”)
  - 7 - miscellaneous (e.g., `man 7 signal` - “What are all the Linux signals”)
- Note, many wrapper functions from the C standard library are named similarly as their system call counterparts. **YOU MUST USE** system calls for your assignment (section 2 ONLY)
  - Read the man page on how to call the right version

# objdump

- Displays information about one or more object files
- usage:
  - `objdump -h <exe>`
  - `objdump -D <object_file>`

```
leaky_prog.o:      file format elf64-x86-64

Disassembly of section .text:

0000000000000000 <main>:
 0:  f3 0f 1e fa                endbr64
 4:  55                        push    %rbp
 5:  48 89 e5                  mov     %rsp,%rbp
 8:  48 83 ec 20               sub     $0x20,%rsp
 c:  89 7d ec                  mov     %edi,-0x14(%rbp)
 f:  48 89 75 e0               mov     %rsi,-0x20(%rbp)
13:  bf 40 00 00 00           mov     $0x40,%edi
18:  e8 00 00 00 00           callq   1d <main+0x1d>
1d:  48 89 45 f8               mov     %rax,-0x8(%rbp)
21:  48 8b 45 f8               mov     -0x8(%rbp),%rax
25:  8b 00                     mov     (%rax),%eax
27:  89 c6                     mov     %eax,%esi
29:  48 8d 3d 00 00 00 00      lea     0x0(%rip),%rdi      # 30 <main+0x30>
30:  b8 00 00 00 00           mov     $0x0,%eax
35:  e8 00 00 00 00           callq   3a <main+0x3a>
3a:  b8 00 00 00 00           mov     $0x0,%eax
3f:  c9                        leaveq  %eax
40:  c3                        retq
```



# nm

- List symbols from object files
- Usage:
  - `nm <exe>`

```
og@ug00 ~/CMPT-379-Labs/Lab-1$ nm leaky_prog
0000000000003dc0 d _DYNAMIC
0000000000003fb0 d _GLOBAL_OFFSET_TABLE_
0000000000002000 R _IO_stdin_used
                                w _ITM_deregisterTMCloneTable
                                w _ITM_registerTMCloneTable
0000000000002154 r __FRAME_END__
0000000000002008 r __GNU_EH_FRAME_HDR
0000000000004010 D __TMC_END__
0000000000004010 B __bss_start
                                w __cxa_finalize@@GLIBC_2.2.5
0000000000004000 D __data_start
0000000000001120 t __do_global_dtors_aux
0000000000003db8 d __do_global_dtors_aux_fini_array_entry
0000000000004008 D __dso_handle
0000000000003db0 d __frame_dummy_init_array_entry
                                w __gmon_start__
0000000000003db8 d __init_array_end
0000000000003db0 d __init_array_start
0000000000001220 T __libc_csu_fini
00000000000011b0 T __libc_csu_init
                                U __libc_start_main@@GLIBC_2.2.5
0000000000004010 D __edata
0000000000004018 B __end
0000000000001228 T __fini
0000000000001000 t __init
0000000000001080 T __start
0000000000004010 b completed.8061
0000000000004000 W data_start
00000000000010b0 t deregister_tm_clones
```

# strace

- Trace system calls and signals
- Usage:
  - `strace [options] command [args]`

[illegible]

# More Linux Tools

- **top** - a tool to show running processes
- **objdump** - display information from object files
- **pstree** - show running processes as a tree
- **ps** - get the list of running processes (-el, -aux)
  - `ps -U CCID -u CCID 1` - get all processes from user `CCID`
  - `ps aux` get all processes on computer!
  - `ps aux | grep <pattern>` returns relevant processes
- **tmux** - terminal multiplexer
  - `Ctrl-b d` to detach, `Ctrl-b %` to split, `Ctrl-b o` next pane
- **[p]kill** - kill a process
  - `pkill -u CCID -U CCID` - kill all processes from user `CCID`
- **watch** - execute a program periodically
  - `watch -n <interval> command`

# Using system calls - error handling

- For most system calls, a return value  $< 0$  indicates an error
- See **man 3 errno** to see all possible errors
- Check the variable **errno** to see what the error is
  - Include the header `#include<errno.h>`
- Use **perror()** to print error detail
  - Include the header `#include<stdio.h>`

# Using system calls

- For this assignment you have to use system calls to implement your shell, here are some listed
  - `chdir()`
  - `getpid()`
  - `fork()`
  - `execve()`
  - `_exit()`
  - `wait()`, `waitpid()`
- Discuss during the next lab
  - `open()`
  - `close()`
  - `dup2()`
  - `pipe()`
  - `kill()`
  - `sigaction()`

```
$ man 2 fork
```

```
$ man 2 execve
```

```
$ man 2 _exit
```

# System call `chdir()`

- Change the current working directory
- Example code - `chdir.c`

```
#include <unistd.h>
```

```
int chdir(const char *path);
```

- Parameter:
  - `path` - which the user want to make the current working directory
- Return Value:
  - `0` - success
  - `-1` - an error occurs and `errno` is set appropriately

# System call `fork()`

- Create a new process
- Example code - fork.c

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
pid_t fork(void);
```

- Return Value:
  - -1 - creation of a child process was unsuccessful
  - 0 - Returned to the newly created child process
  - >0 - Returned to parent or caller. The value contains process ID of newly created child process

# System call `getpid()`, `getppid()`

- Get the process ID
- Example code - `getpid.c`

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
pid_t getpid(void);
```

- Return Value:
  - The process ID of the calling process

```
pid_t getppid(void);
```

- Return Value:
  - The process ID of the parent process



# System call `execve()`

- Execute a program and replace the current process image
- Example code - `execve.c`

```
#include <unistd.h>
```

```
int execve(const char *path, char *const argv[], char *const envp[]);
```

- Parameters:
  - `path` - the path of the file being executed
  - `argv` - null terminated array of the arguments for the program being executed
  - `envp` - array of strings, conventionally of the form **key=value**
- Return Value:
  - No return - success
  - -1 - an error occurs and **errno** is set appropriately

# System call `_exit()`

- Terminate process and return status to the parent
- Not the same as `exit()` **DO NOT USE** `exit()` on the assignment as it is a library function
- Example codes - `exit1.c` and `exit2.c`

```
#include <unistd.h>
```

```
int _exit(int status);
```

- Parameter:
  - Status - value returned to the parent process

# System call `wait()`

- Wait until one of its children terminates
- Example code - `wait.c` and `waitpid.c`

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

```
pid_t wait(int *wstatus);
```

- Parameter:
  - Status - value returned to the parent process
- Return Value:
  - Terminated process ID - success
  - 0 or -1 - error

# System call `waitpid()`

- Wait for a specific process ID to change state
- Example code - waitpid.c

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

```
pid_t waitpid(pid_t pid, int *wstatus, int options);
```

- Parameter:
  - `pid` - The process id to wait on
  - `wstatus` - Value returned to the parent process (often `NULL`)
  - `options` - Ored constants to change the behaviour of `waitpid` (often `0`)
- Return Value:
  - Terminated process ID - success
  - `0` or `-1`

\$ man 2 times



## System call `times()`

- Get the process times
- Example code - times.c

```
#include <sys/times.h>
```

```
struct tms {  
    clock_t tms_utime; /* user time */  
    clock_t tms_stime; /* system time */  
    clock_t tms_cutime; /* user time of children */  
    clock_t tms_cstime; /* system time of children */  
};
```

```
clock_t times(struct tms *buf);
```

- Return Value:
  - The number of clock ticks that have elapsed since the past (> 0)
  - -1 - error has occurred and errno is set

# System call `open()`

- Open a file

```
#include <sys/stat.h>
```

```
#include <fcntl.h>
```

```
int open(const char *path, int oflags, mode_t mode);
```

- Parameter:
  - `oflags` - `O_RDONLY`, `O_WRONLY`, `O_RDWR`, `O_APPEND`, `O_CREAT`, etc.
  - `mode` - `S_IRUSR`, `S_IWUSR`, `S_IXUSR`, etc.
- Return Value:
  - `-1` - error
  - Others - file descriptor

# System call `close()`

- Close a file

```
#include <unistd.h>
```

```
int close(int fildes);
```

- Parameter:
  - `fildes` - The file descriptor to be closed
- Return Value:
  - `-1` - error
  - `0` - success

# System call `write()`

- Write to a file descriptor
- Example code - write.c

```
#include <unistd.h>
```

```
ssize_t write(int fd, const void* buf, size_t count);
```

- Parameter:
  - fd - The file descriptor to be written to
  - buf - The buffer to write to the fd
  - count - The number of bytes to write to the fd
- Return Value:
  - -1 - error
  - 0 - success



# System call `read()`

- Read from a file descriptor
- Example code - `read.c`

```
#include <unistd.h>
```

```
ssize_t read(int fd, const void* buf, size_t count);
```

- Parameter:
  - `fd` - The file descriptor to be read from
  - `buf` - The buffer to read the contents in `fd` into
  - `count` - The number of bytes to read from the `fd`
- Return Value:
  - `-1` - error
  - `0` - success

# System call `dup2()`

- Duplicates one file descriptor, making them aliases, and then deleting the old file descriptor
- Example code - dup2.c

```
#include <unistd.h>
```

```
int dup2(int fildes, int fildes2);
```

- Parameter:
  - `fildes` - source file descriptor
  - `fildes2` - target file descriptor
- Return Value:
  - `<0` - error
  - Others - second file descriptor

# System call `pipe()`

- Creates a unidirectional data channel for interprocess communication (IPC)
- Example code - pipe.c

```
int pipe(int pipefd[2]);
```

- Parameter:
  - `pipefd` - two file descriptors, read/write ends of the pipe
- Return Value:
  - `-1` - error
  - `0` - success

# System call `kill()`

- Send signal to a process
- Example code - kill.c

```
#include <sys/types.h>
```

```
#include <signal.h>
```

```
int kill(pid_t pid, int sig);
```

- Parameter:
  - `pid` - target process
  - `sig` - signal want to send
- Return Value:
  - `0` - success
  - `-1` - error

# Question from class

What if the parent of a process is killed?

systemd(1) - A - B - C

`pstree -s -p <pid>`

A gets killed... B is orphan to systemd(1)

systemd(1) - B - C

`ps -f <pid>`

B gets killed... C is orphan to systemd(1), B is a zombie

systemd(1) - C

systemd(1) - A - B(Z)