

CMPUT-379 Lab 2

Akemi Izuko, Armaan Katyal, Ellis McDougald, Han Yang, Patrick Zijlstra, Shasta Johnsen-Sollos, Steven Oufan Hai, Zhaoyu Li

Today's lab

- Pipes
- Shared memory
 - Producer/consumer example
- Signals
- Using system calls in c programs
- Environment variables in Linux
- Review of assignment 1

Pipes

- Interprocess Communication (IPC)
 - Half-duplex (unidirectional)
 - Inherited by child processes
-
- For reference: Ch. 15 APUE.3e


System call `pipe()`

- Creates a unidirectional data channel for interprocess communication (IPC)
- Example code - pipe.c


```
int pipe(int pipefd[2]);
```

- Parameter:
 - `pipefd` - gets populated with the file descriptors of the r/w ends
- Return Value:
 - `-1` - error, sets `errno`
 - `0` - success

Pipe - talk to self



Pipe - talk to child process



Pipe - redirect child's output to parent

```
int main (int argc, char *argv[]) {
    char buf[MAXBUF];
    int n, status, fd[2];
    pid_t pid;
    if (pipe(fd) < 0) perror("pipe error!");
    if ((pid = fork()) < 0) perror("fork error!");
    if (pid == 0) {
        close(fd[0]);                // child won't read
        dup2(fd[1], STDOUT_FILENO); // stdout = fd[1]
        close(fd[1]);                // stdout is still open
        if (execl("/usr/bin/w", "w", (char *) 0) < 0) perror("execl error!"); //
    } else {
        close(fd[1]);                // parent won't write
        while ((n = read(fd[0], buf, MAXBUF)) > 0)
            write(STDOUT_FILENO, buf, n);
        close(fd[0]);
        wait(&status);
    }
```

\$ w | wc -w

```
int main (int argc, char *argv[]) {
    int fd[2]; pid_t pid;

    if (pipe(fd) < 0) perror("pipe error!");
    if ((pid = fork()) < 0) perror("fork error!");
    if (pid == 0) {
        close(fd[1]);                // child won't write
        dup2(fd[0], STDIN_FILENO);    // stdin = fd[0]
        close(fd[0]);                // stdin is still open
        if (execl("/usr/bin/wc", "wc", "-w", (char *) 0) < 0)
            perror("execl error!");
    } else {
        close(fd[0]);                // parent won't read
        dup2(fd[1], STDOUT_FILENO);   // stdout = fd[1]
        close(fd[1]);                // stdout is still open
        if (execl("/usr/bin/w", "w", (char *) 0) < 0)
            perror("execl error!");
    }
    return 0;
}
```



Interprocess Communication (IPC)

- A process is **Independent** if it does not share data
- A process is **Cooperating** if it can affect or be affected by the other processes executing in the system




Information Sharing

- Several applications may be interested in the same information
- We must provide an environment to allow concurrent access
- Cooperating processes require an IPC mechanism to allow for data exchanges
- Two models
 - Shared Memory
 - Message Passing
- Several ways
 - mmap (shared memory)
 - System V message passing
 - Send signals
 - Pipes
 - Sockets




Shared Memory

- A shared region of memory is established
- Processes can exchange information by reading and writing to this region



Shared Memory

- Shared memory is typically faster than message passing due to less system calls and kernel intervention
- System calls are only required to establish shared memory regions
- Once established, all accesses are treated as routine memory accesses



Functions for shared memory

```
#include <sys/mman.h>
#include <sys/stat.h>      /* For mode constants */
#include <fcntl.h>         /* For O_* constants */
```

```
int shm_open(const char *name, int oflag, mode_t mode);
int shm_unlink(const char *name);
```

```
#include <sys/mman.h>
void *mmap(void *addr, size_t length, int prot, int flags, int fd, off_t offset);
int munmap(void *addr, size_t length);
```

LINK WITH -lrt!!!!!!! (real-time library)

Shared Memory Example - producer.c

```
int main() {
    const int SIZE = 4096;
    const char* name = "/PC";

    const char* message_0 = "Hello";
    const char* message_1 = "World!";

    int shm_fd;
    void* ptr;

    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);
    ftruncate(shm_fd, SIZE);
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    sprintf(ptr, "%s", message_0);
    ptr += strlen(message_0);
    sprintf(ptr, "%s", message_1);
    ptr += strlen(message_1);

    return 0;
}
```

/* size (B) of shared memory object */
/* name of the shared memory object */

/* shared memory file descriptor */
/* pointer to shared memory object */

/* create the shared memory object */
/* configure size of the shared memory object */
/* memory map the shared memory object */

/* write to the shared memory object */

Shared Memory Example - consumer.c

```
int main() {
    const int SIZE = 4096;                /* size (B) of shared memory object */
    const char* name = "/PC";            /* name of the shared memory object */

    int shm_fd;                          /* shared memory file descriptor */
    void* ptr;                          /* pointer to shared memory object */

    shm_fd = shm_open(name, O_RDONLY, 0666); /* open the shared memory object */

    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0); /* memory map the shared memory object */

    printf("%s", (char*)ptr);            /* read from the shared memory object */
    shm_unlink(name);                   /* remove the shared memory object */

    return 0;
}
```

Shared Memory Concurrency

- We do not want two processes to access the shared memory at the same time
- Semaphores are a signaling mechanism that act similar to a lock




Shared Memory Concurrency

Good Case

| Thread 1 | Thread 2 |
|---------------------|---------------------|
| ----- | |
| 1. Load to Reg. (0) | |
| 2. Incr. Reg. (1) | doing something |
| 3. Store Reg. (1) | else |
| | 1. Load to Reg. (1) |
| doing something | 2. Incr. Reg. (2) |
| else | 3. Store Reg. (2) |


Result: counter = 2, as expected!



Shared Memory Concurrency

| Thread 1 | Thread 2 |
|---------------------|---------------------|
| 1. Load to Reg. (0) | |
| 2. Incr. Reg. (1) | 1. Load to Reg. (0) |
| 3. Store Reg. (1) | 2. Incr. Reg. (1) |
| | 3. Store Reg. (1) |

Result: counter = 1



Shared Memory Concurrency

Solutions:

- Using **signals** to coordinates different processes
- FYI: Using **Semaphore** to restrict the maximum number of processes accessing the shared memory at the same time
- FYI: Using **inter-process mutex** to lock shared memory when a process is accessing
 - Mutex usually use for multithreading
 - To use mutex in multiprocessing, the mutex need to be allocated in shared memory

Signal

- Software interrupt
- Use **signal()** or **sigaction()** to install signal handlers
- **int kill(pid_t pid, int sig)** sends a signal
- For reference: Ch. 10 APUE.3e

Signal Handling with **signal()**

```
#include <signal.h>
```

```
sighandler_t signal(int signum, sighandler_t handler);
```

```
// Sets the disposition of signal signum to handler
```

- Parameters:
 - `signum`: Specifies the signal
 - `handler`: A function that takes a single `int` and returns nothing
- Return Value:
 - On success: The previous value of the signal handler
 - On failure: `SIG_ERR`

Portability Problem: Behavior across UNIX versions may vary

Signal handling with **signal()**

```
int i;
void quit(int signum) {
    fprintf(stderr, "\nInterrupt (code= %d, i= %d)\n", signum, i);
}
int main () {
    if(signal(SIGQUIT, quit) == SIG_ERR)
        perror("can't catch SIGQUIT");
    for (i= 0; 1; i++) {
        usleep(1000);
        if (i % 100 == 0) putc('.', stderr);
    }
    return(0);
} // signal/signal2.c
```

```
$ ./signal2
```

```
.....^\  
Interrupt (code= 3, i= 752)
```

```
.....^\  
Interrupt (code= 3, i= 1416)
```

```
.....^\  
Interrupt (code= 3, i= 2336)
```

```
.....^C
```

Signal Handling with **sigaction()**

```
#include <signal.h>
```

```
int sigaction(int signum, const struct sigaction *act, struct  
sigaction *oldact);
```

```
// Changes the action taken by a process on a receipt of a  
specific signal
```

- Parameters:
 - **signum**: Specifies the signal - Cannot be **SIGKILL** or **SIGSTOP**
 - **act**: If non-NULL, this is the new action for signal **signum**
 - **oldact**: If non-NULL, the previous action is saved in **oldact**
- Return Value:
 - 0 - success
 - -1 - error

sigaction struct

```
struct sigaction {  
    void (*sa_handler)(int);    // signal handler function  
                                // or SIG_IGN (ignore)  
                                // or SIG_DFL (default)  
    sigset_t sa_mask;           // additional signals to block  
    int sa_flags;               // signal options  
    void (*sa_handler)(int, siginfo_t *, void *) // alternate handler  
}
```


siginfo struct

- Set SA_SIGINFO to use alternate handler to get information about signal

```
struct siginfo{  
    int si_signo;           // signal number  
    int si_errno;          // errno value  
    int si_code;           // more specific signal category  
    pid_t si_pid;          // sending process id  
    uid_t si_uid;          // sending process user id  
    void *si_addr;         // address that caused signal  
    ...  
}
```

Signal handling with **sigaction()**

```
$ ./signal1
```

```
^CCaught signal!
```

```
^CCaught signal!
```

```
^CCaught signal!
```

```
// signal/signal1.c
```

```
void signal_callback_handler(int signum) {  
    printf("Caught signal!\n");  
}
```

The `sigset_t` data type is an array of booleans representing a set of signals

`sigset_t s` 

```
int main() {  
    struct sigaction sa;  
    sa.sa_flags = 0;  
    sigemptyset(&sa.sa_mask); // clears  
    sa.sa_handler = signal_callback_handler;  
    sigaction(SIGINT, &sa, NULL);  
    // sigaction(SIGTSTP, &sa, NULL);  
    while (1) {}  
}
```

`sigemptyset(&s);`

Using system calls

- For this assignment you have to use system calls to implement your shell.

Continuing from last week, here are a few more:

- `open()`
- `close()`
- `write()`
- `read()`
- `dup2()`
- `kill()`
- `nice()`
- `mmap()`
- `munmap()`

- Discussed earlier:

- `pipe()`
- `signal()`
- `sigaction()`

System call `open()`

- Open a file

```
#include <sys/stat.h>
```

```
#include <fcntl.h>
```

```
int open(const char *path, int oflags, mode_t mode);
```

- Parameter:
 - `oflags` - `O_RDONLY`, `O_WRONLY`, `O_RDWR`, `O_APPEND`, `O_CREAT`, etc.
 - `mode` - `S_IRUSR`, `S_IWUSR`, `S_IXUSR`, etc.
- Return Value:
 - `-1` - error
 - Others - file descriptor

System call `close()`

- Close a file

```
#include <unistd.h>
```

```
int close(int fildes);
```

- Parameter:
 - `fildes` - The file descriptor to be closed
- Return Value:
 - `-1` - error
 - `0` - success

System call `write()`

- Write to a file descriptor
- Example code - write.c

```
#include <unistd.h>
```

```
ssize_t write(int fd, const void* buf, size_t count);
```

- Parameter:
 - fd - The file descriptor to be written to
 - buf - The buffer to write to the fd
 - count - The number of bytes to write to the fd
- Return Value:
 - -1 - error
 - Number of bytes written - success

System call `read()`

- Read from a file descriptor
- Example code - `read.c`

```
#include <unistd.h>
```

```
ssize_t read(int fd, const void* buf, size_t count);
```

- Parameter:
 - `fd` - The file descriptor to be read from
 - `buf` - The buffer to read the contents in `fd` into
 - `count` - The number of bytes to read from the `fd`
- Return Value:
 - `-1` - error
 - Number of bytes read - success (`0` on EOF)

System call `dup2()`

- Duplicates one file descriptor, making them aliases, and then deleting the old file descriptor
- Example code - dup2.c

```
#include <unistd.h>
```

```
int dup2(int fildes, int fildes2);
```

- Parameter:
 - `fildes` - source file descriptor
 - `fildes2` - target file descriptor
- Return Value:
 - `<0` - error
 - Others - second file descriptor

System call `kill()`

- Send signal to a process
- Example code - kill.c

```
#include <sys/types.h>
```

```
#include <signal.h>
```

```
int kill(pid_t pid, int sig);
```

- Parameter:
 - `pid` - target process
 - `sig` - signal want to send
- Return Value:
 - `0` - success
 - `-1` - error

System call `nice()`

- Change process priority (higher nice value → lower priority)

```
#include <unistd.h>
```

```
int nice(int incr);
```

- Parameter:
 - `incr` - amount to add to current nice value
- Return Value:
 - new nice value - success
 - -1 - error

System call `mmap()`

- Map files into memory

```
#include <sys/mman.h>
```

```
int mmap(void *addr, size_t len, int prot, int flag, int fd, off_t off);
```

- Parameter:
 - `addr` - start address of mapping
 - `len` - length of mapping
 - `prot` - desired memory protection of mapping
 - `flag` - behaviour when updating mapping
 - `fd` - file to get content from
 - `off` - offset to start from within file
- Return Value:
 - start address of mapped region - success
 - `MAP_FAILED` - error

System call `munmap()`

- Unmap files from memory

```
#include <sys/mman.h>
```

```
int munmap(void *addr, size_t len);
```

- Parameter:
 - `addr` - start address of mapping
 - `len` - length of mapping
- Return Value:
 - `0` - success
 - `-1` - error

Environment Variables in Linux

- Named strings that are accessible by all applications (“ultra global” variables)
- Typically used for defining environment specific configurations

`NAME=value`

- Common env variables that are set automatically:
 - `PATH`: list of paths to search for executable files
 - `HOME`: home directory
 - etc.
- For reference: Section 7.9 of APUE.3e

Modifying Environment Variables

- Places the provided string of form name=value into the environment list
- Always overwrites if variable exists: if name already exists, replaces with new value

```
#include <stdlib.h>
```

```
int putenv(char *str);
```

- Parameter:
 - str - string to place into environment list
- Return Value:
 - 0 - success
 - !0 - error

Modifying Environment Variables

- Sets the value for an environment variable
- Can overwrite existing variables if desired

```
#include <stdlib.h>
```

```
int setenv(const char *name, const char *value, int rewrite);
```

- Parameter:
 - name - name of environment variable
 - value - value to set
 - rewrite - whether to overwrite an existing definition
- Return Value:
 - 0 - success
 - -1 - error

Assignment 1

- Build an interactive command line shell (dragonshell)
- Gain experience with system calls for process management and IPC
- Remember: for core features, only use system call functions from Section 2 of the man pages
- Familiarize yourself with the rubric on eclass

- Due date:
 - **October 1, 2025 at 5:59 pm**

Assignment 1 - high level approach

Make it easy for yourself! Do one step at a time, don't try to do pieces end-to-end

1. Read line of input into array of strings
2. Construct command struct
3. Validate command struct for inconsistency
4. Run command struct in one of 3 cases: (builtin, pipe, background + redirect)

```
typedef struct Cmd {  
    char* cmd;  
    char** args;  
    int read_from; // fd for read  
    int write_to; // fd for write  
    bool is_background;  
    bool is_builtin;  
    struct Cmd* pipe_to;  
} Cmd;
```

Quiz

Write a program that takes in a single argument, then prints out the unique files in that command sorted. You must use pipe(2) twice!

Similar to running the following in bash... where \$1 is the first argument in argv

```
cat $1 | sort | uniq
```